

Multicollinearity: Question

STA 506 2.0 Linear Regression Analysis

Thiyanga S. Talagala

Data

```
library(tidyverse)
realestate <- read.csv("real-estate.csv")
head(realestate)
```

	ID	Price	Sqft	Bedroom	Bathroom	Airconditioning	Garage	Pool	YearBuild	Quality
1	1	360000	3032	4	4	1	2	0	1972	2
2	2	340000	2058	4	2	1	2	0	1976	2
3	3	250000	1780	4	3	1	2	0	1980	2
4	4	205500	1638	4	2	1	2	0	1963	2
5	5	275500	2196	4	3	1	2	0	1968	2
6	6	248000	1966	4	3	1	5	1	1972	2

	Lot	AdjHighway
1	22221	0
2	22912	0
3	21345	0
4	17342	0
5	21786	0
6	18902	0

```
glimpse(realestate)
```

```
Rows: 522
Columns: 12
$ ID           <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, ...
$ Price        <int> 360000, 340000, 250000, 205500, 275500, 248000, 229...
$ Sqft         <int> 3032, 2058, 1780, 1638, 2196, 1966, 2216, 1597, 162...
$ Bedroom      <int> 4, 4, 4, 4, 4, 4, 3, 2, 3, 3, 7, 3, 5, 5, 3, 5, 2, ...
$ Bathroom     <int> 4, 2, 3, 2, 3, 3, 2, 1, 2, 3, 5, 4, 4, 4, 3, 5, 2, ...
$ Airconditioning <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, ...
$ Garage       <int> 2, 2, 2, 2, 2, 5, 2, 1, 2, 1, 2, 3, 3, 2, 2, 2, 2, ...
$ Pool         <int> 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, ...
$ YearBuild    <int> 1972, 1976, 1980, 1963, 1968, 1972, 1972, 1955, 197...
$ Quality      <int> 2, 2, 2, 2, 2, 2, 2, 2, 3, 3, 3, 3, 1, 1, 1, 2, 2, ...
$ Lot         <int> 22221, 22912, 21345, 17342, 21786, 18902, 18639, 22...
$ AdjHighway   <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
```

Q1: Identify qualitative and quantitative variables.

ID	Price	Sqft	Bedroom		
Min. : 1.0	Min. : 84000	Min. : 980	Min. : 0.000		
1st Qu.: 131.2	1st Qu.: 180000	1st Qu.: 1701	1st Qu.: 3.000		
Median : 261.5	Median : 229900	Median : 2061	Median : 3.000		
Mean : 261.5	Mean : 277894	Mean : 2261	Mean : 3.471		
3rd Qu.: 391.8	3rd Qu.: 335000	3rd Qu.: 2636	3rd Qu.: 4.000		
Max. : 522.0	Max. : 920000	Max. : 5032	Max. : 7.000		
Bathroom	Airconditioning	Garage	Pool	YearBuild	Quality
Min. : 0.000	0: 88	Min. : 0.0	0: 486	Min. : 1885	1: 68
1st Qu.: 2.000	1: 434	1st Qu.: 2.0	1: 36	1st Qu.: 1956	2: 290
Median : 3.000		Median : 2.0		Median : 1966	3: 164
Mean : 2.642		Mean : 2.1		Mean : 1967	
3rd Qu.: 3.000		3rd Qu.: 2.0		3rd Qu.: 1981	
Max. : 7.000		Max. : 7.0		Max. : 1998	
Lot	AdjHighway				
Min. : 4560	0: 511				
1st Qu.: 17205	1: 11				
Median : 22200					
Mean : 24370					
3rd Qu.: 26787					
Max. : 86830					

Q2: What is wrong with the following graph?

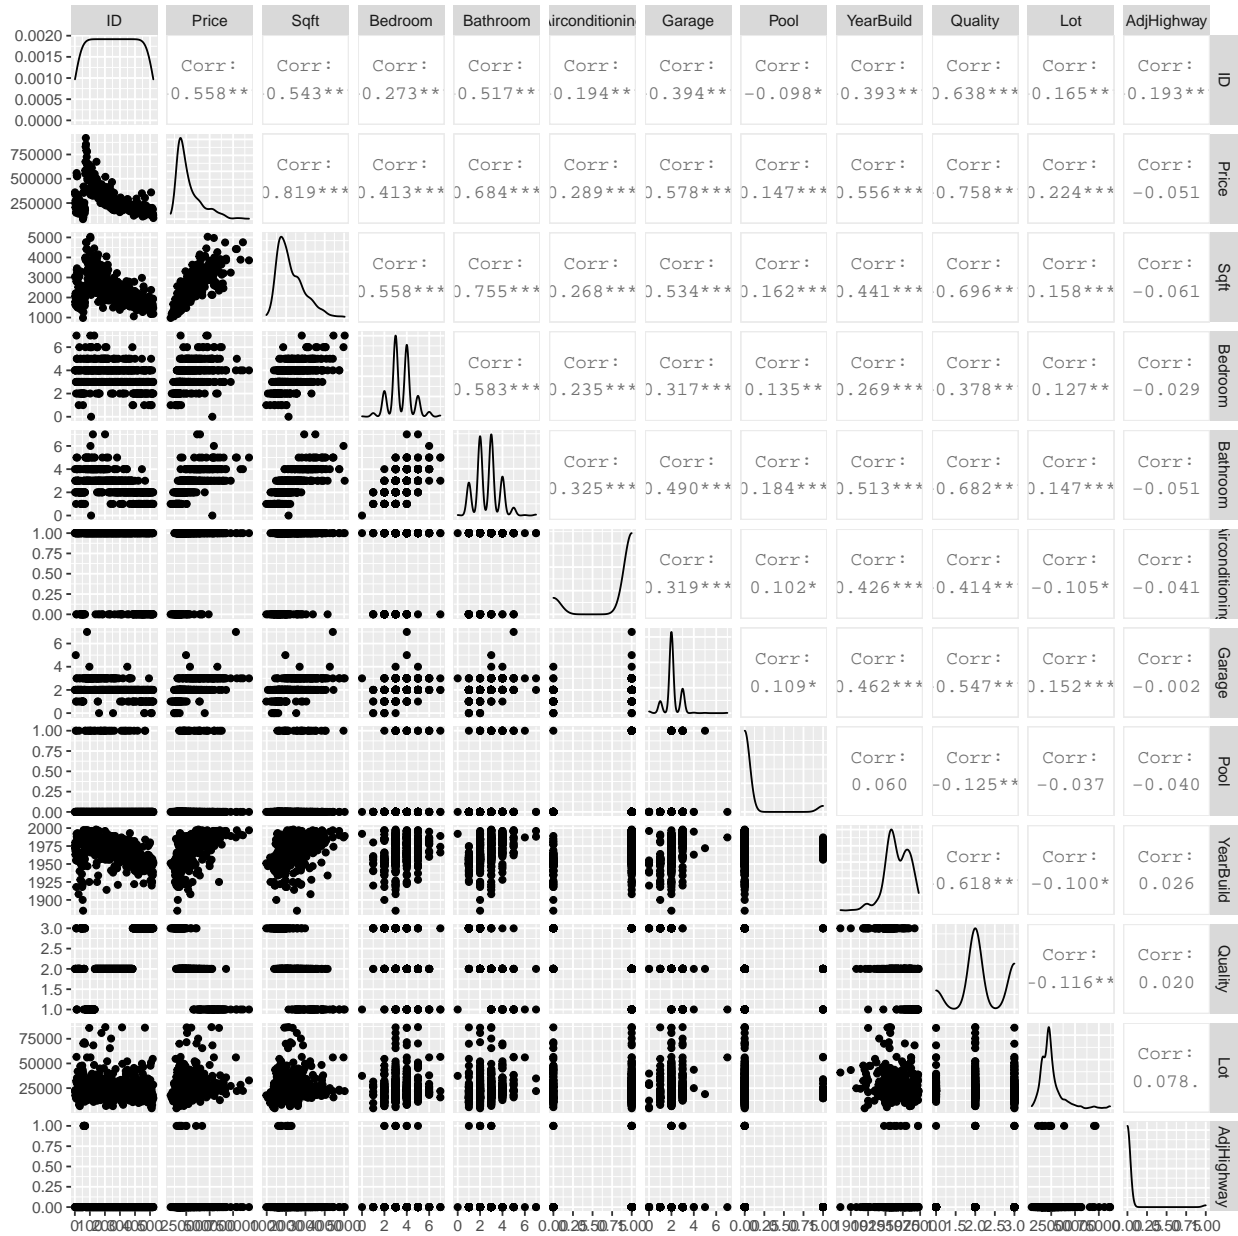


Figure 1: Pairwise correlation plot

Q3: Figure 1 is modified as follows. Now what can you say about the graph?

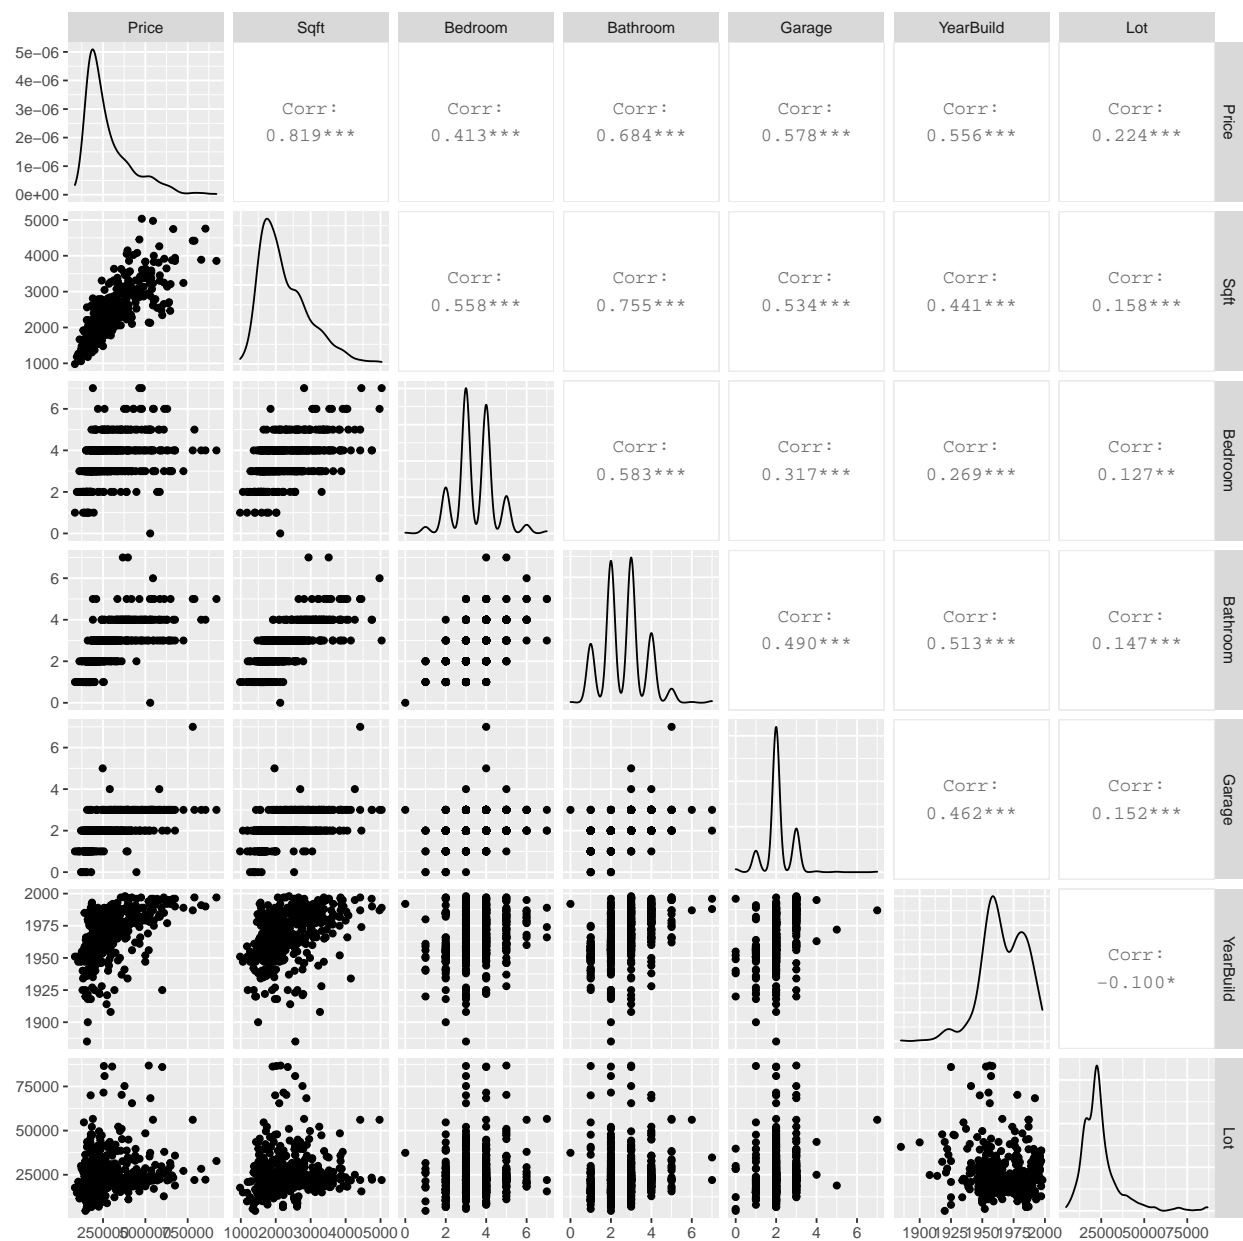


Figure 2: Pairwise correlation plot

Q4: What can you say about the source of multicollinearity in these data?

```
realestate$Airconditioning <- factor(realestate$Airconditioning)
realestate$Pool <- factor(realestate$Pool)
realestate$AdjHighway <- factor(realestate$AdjHighway)
realestate$Quality <- factor(realestate$Quality)
realestate <- realestate[, -1]
summary(realestate)
```

Price	Sqft	Bedroom	Bathroom
Min. : 84000	Min. : 980	Min. : 0.000	Min. : 0.000
1st Qu.: 180000	1st Qu.: 1701	1st Qu.: 3.000	1st Qu.: 2.000
Median : 229900	Median : 2061	Median : 3.000	Median : 3.000
Mean : 277894	Mean : 2261	Mean : 3.471	Mean : 2.642
3rd Qu.: 335000	3rd Qu.: 2636	3rd Qu.: 4.000	3rd Qu.: 3.000
Max. : 920000	Max. : 5032	Max. : 7.000	Max. : 7.000

Airconditioning	Garage	Pool	YearBuild	Quality	Lot
0: 88	Min. : 0.0	0: 486	Min. : 1885	1: 68	Min. : 4560
1: 434	1st Qu.: 2.0	1: 36	1st Qu.: 1956	2: 290	1st Qu.: 17205
	Median : 2.0		Median : 1966	3: 164	Median : 22200
	Mean : 2.1		Mean : 1967		Mean : 24370
	3rd Qu.: 2.0		3rd Qu.: 1981		3rd Qu.: 26787
	Max. : 7.0		Max. : 1998		Max. : 86830

AdjHighway

0: 511

1: 11

```
model <- lm(Price ~ . , data=realestate)
summary(model)
```

Call:

```
lm(formula = Price ~ . , data = realestate)
```

Residuals:

Min	1Q	Median	3Q	Max
-204865	-28010	-4973	21315	298892

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-2.358e+06	3.991e+05	-5.909	6.29e-09 ***
Sqft	8.700e+01	6.570e+00	13.242	< 2e-16 ***
Bedroom	-5.125e+03	3.275e+03	-1.565	0.1182
Bathroom	8.127e+03	4.288e+03	1.895	0.0586 .
Airconditioning1	4.851e+03	8.086e+03	0.600	0.5488
Garage	1.089e+04	5.060e+03	2.152	0.0319 *
Pool1	1.014e+04	1.040e+04	0.975	0.3303
YearBuild	1.269e+03	2.024e+02	6.272	7.60e-10 ***

Quality2	-1.430e+05	1.021e+04	-14.007	< 2e-16	***
Quality3	-1.484e+05	1.404e+04	-10.564	< 2e-16	***
Lot	1.556e+00	2.363e-01	6.587	1.12e-10	***
AdjHighway1	-2.737e+04	1.810e+04	-1.512	0.1311	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 58770 on 510 degrees of freedom
Multiple R-squared: 0.8223, Adjusted R-squared: 0.8184
F-statistic: 214.5 on 11 and 510 DF, p-value: < 2.2e-16

```
library(car)
car::vif(model)
```

	GVIF	Df	GVIF ^{1/(2*Df)}
Sqft	3.292569	1	1.814544
Bedroom	1.664845	1	1.290289
Bathroom	3.141563	1	1.772445
Airconditioning	1.385038	1	1.176876
Garage	1.651938	1	1.285277
Pool	1.050442	1	1.024911
YearBuild	1.922344	1	1.386486
Quality	3.322305	2	1.350081
Lot	1.150133	1	1.072443
AdjHighway	1.021444	1	1.010665

Note:

1.1: note

Multicollinearity occurs if we do not treat this appropriately.

1.2 Model with only quantitative variables: VIF

```
summary(Duncan)
```

	type	income	education	prestige
bc	:21	Min. : 7.00	Min. : 7.00	Min. : 3.00
prof	:18	1st Qu.:21.00	1st Qu.: 26.00	1st Qu.:16.00
wc	: 6	Median :42.00	Median : 45.00	Median :41.00
		Mean :41.87	Mean : 52.56	Mean :47.69
		3rd Qu.:64.00	3rd Qu.: 84.00	3rd Qu.:81.00
		Max. :81.00	Max. :100.00	Max. :97.00

```
m1 <- lm(prestige ~ income + education, data=Duncan)
vif(m1)
```

	income	education
	2.1049	2.1049

1.3 Model with only quantitative variables and qualitative variables: Generalized variance-inflation factors

```
m2 <- lm(prestige ~ income + education + type, data=Duncan)
vif(m2)
```

	GVIF	Df	GVIF^(1/(2*Df))
income	2.209178	1	1.486330
education	5.297584	1	2.301648
type	5.098592	2	1.502666

Q5: Write down the estimated model.

```
summary(model)
```

Call:

```
lm(formula = Price ~ ., data = realestate)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-204865	-28010	-4973	21315	298892

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-2.358e+06	3.991e+05	-5.909	6.29e-09	***
Sqft	8.700e+01	6.570e+00	13.242	< 2e-16	***
Bedroom	-5.125e+03	3.275e+03	-1.565	0.1182	
Bathroom	8.127e+03	4.288e+03	1.895	0.0586	.
Airconditioning1	4.851e+03	8.086e+03	0.600	0.5488	
Garage	1.089e+04	5.060e+03	2.152	0.0319	*
Pool1	1.014e+04	1.040e+04	0.975	0.3303	
YearBuild	1.269e+03	2.024e+02	6.272	7.60e-10	***
Quality2	-1.430e+05	1.021e+04	-14.007	< 2e-16	***
Quality3	-1.484e+05	1.404e+04	-10.564	< 2e-16	***
Lot	1.556e+00	2.363e-01	6.587	1.12e-10	***
AdjHighway1	-2.737e+04	1.810e+04	-1.512	0.1311	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 58770 on 510 degrees of freedom

Multiple R-squared: 0.8223, Adjusted R-squared: 0.8184

F-statistic: 214.5 on 11 and 510 DF, p-value: < 2.2e-16